



Kanten zwischen den Hexas für TGGs und "halbe" Hexas.
Schön eigentlich.

A Precedence-Driven Approach for Concurrent Model Synchronization Scenarios using Triple Graph Grammars

Anonymous Author(s)

Abstract

Concurrent model synchronization is the task of restoring consistency between two correlated models after they have been changed concurrently and independently. To determine whether such concurrent model changes conflict with each other and to resolve these conflicts taking domain- or user-specific preferences into account is highly challenging. In this paper, we present a framework for concurrent model synchronization algorithms based on Triple Graph Grammars (TGGs). TGGs specify the consistency of correlated models using grammar rules; these rules can be used to derive different consistency restoration operations. Using TGGs, we infer a causal dependency relation for model elements that enables us to detect conflicts non-invasively. Different kinds of conflicts are detected first and resolved by the subsequent conflict resolution process. Users configure the overall synchronization process by orchestrating the application of consistency restoration fragments according to several conflict resolution strategies to achieve individual synchronization goals. As proof of concept, we have implemented this framework in the model transformation tool eMoflon. Our initial evaluation shows that the runtime of our presented approach scales with the size of model changes and conflicts, rather than model size.

CCS Concepts: • Software and its engineering → Model-driven software engineering; Consistency; Software evolution.

Keywords: bidirectional transformation (bx), concurrent model synchronization, triple graph grammars

ACM Reference Format:

Anonymous Author(s). 2020. A Precedence-Driven Approach for Concurrent Model Synchronization Scenarios using Triple Graph Grammars. In *Proceedings of ACM SIGPLAN International Conference*

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

SLE '20, November 15–20, 2020, Chicago, IL
© 2020 Association for Computing Machinery.
ACM ISBN 978-x-xxxx-xxxx-x/YY/MM...\$15.00
<https://doi.org/10.1145/nnnnnnn.nnnnnnn>

on *Software Language Engineering (SLE '20)*. ACM, New York, NY, USA, 17 pages. <https://doi.org/10.1145/nnnnnnn.nnnnnnn>

1 Introduction

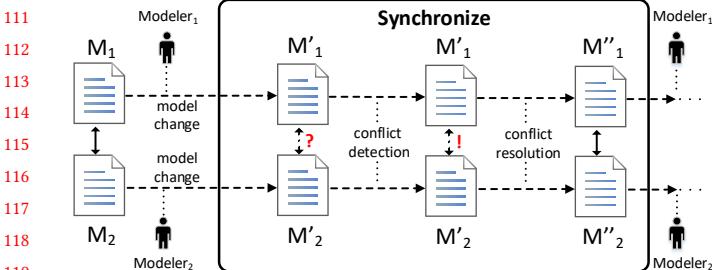
Model-driven engineering [4] has proven to be an effective means to tackle the challenges that accompany the development of modern software systems, which are getting increasingly complex and distributed in nature. Often more than one model is needed to describe the developed software system from different, but overlapping perspectives. Keeping these models and various types of traceability relationships between them in a consistent state is a challenging task, often called model synchronization.

Model synchronization becomes especially challenging when multiple correlated models are changed concurrently. In such cases, not all changes can always be propagated between models as some may contradict each other and thus, are in conflict. This is the case, for example, when a change in one model leads to the deletion of elements in the other whose existence is the prerequisite for changes performed in that second model by another user. Yet, even for model changes that are not in conflict, there may be multiple ways to propagate them between models.

For a modern concurrent synchronization approach, it is of paramount importance to identify synchronization conflicts reliably and give modelers the ability to orchestrate model synchronization processes for guiding the process in accordance with their goals. Figure 1 gives an overview of a concurrent model synchronization process. Consistent interrelated models M_1 and M_2 are given and changed concurrently. The synchronization process identifies all conflicts between these changes and runs a conflict resolution process. The expected synchronization result is a consistent pair of models M'_1 and M'_2 that contains all conflict-free changes.

Some approaches provide solely one hard-wired solution of a conflict-detection and resolution strategy (from a universe of many different options), others come without any formal guarantees for their synchronization results or have an exponential runtime behavior w.r.t. the size of the processed models (cf. Section 2).

The contribution of this paper is a framework that simplifies the implementation (orchestration of a family) of concurrent synchronization algorithms with the following properties:

**Figure 1.** Synchronization Process

- These algorithms are derived from a declarative rule-based formal specification of a model consistency relation in the form of so-called Triple Graph Grammars (TGGs).
- They come with a number of predefined but extensible conflict-detection and -resolution as well as consistency restoration strategies.
- Formal properties can be shown using state-of-the-art category- and graph-transformation-based proof techniques (e.g., [11, 14, 21, 30])
- The intended scope of the effects and potential conflicts of model changes are identified using a TGG-based causal dependency relation for model elements.
- Scalability with the size of processed model changes is guaranteed by limiting the effects of model updates to causally dependent areas in the regarded models and relying on incremental graph pattern matching techniques.

In Section 2, we give an overview of the state-of-the-art w.r.t. concurrent model synchronization approaches. Section 3 recalls various concepts related to TGGs. Sections 4 and 5 present our concurrent synchronization framework. The former explains in detail how conflicts are detected, while the latter presents strategies to resolve them and restore consistency. In Section 6, we introduce our implementation. Based on this, we evaluate our approach w.r.t. scalability in Section 7. Section 8 sums up our contribution and discusses future work.

2 Related Work

In this section, we discuss the current state-of-the-art in the field of concurrent model synchronization. Although we do not claim completeness of this survey, we are not aware of further works that differ fundamentally from the ones discussed here. The majority of the approaches in this field can be categorized into *propagation-based*, *constraint-based*, and *search-based* approaches. The approaches considered support *state-based* and *delta-based* model changes; however, they cannot be clearly clustered as some approaches abstract from the way how model changes are described and,

therefore, support both state- and delta-based definitions of model changes.

State-based approaches [37–39] hold copies of all models in order to calculate differences, which is not only memory consuming, but also scales with the size of the involved models. In contrast, *delta-based* approaches [16, 17, 19, 35] operate on model changes, which may, e.g., be detected by an *incremental pattern matcher*. In general, delta-based approaches tend to scale better in scenarios with frequent changes, but require more bookkeeping, which may have negative effects to the memory consumption.

Propagation-based approaches to concurrent model synchronization use sequential synchronization steps to propagate the changes from one model to the other one followed by a propagation step in the opposite direction. All propagation-based approaches have severe drawbacks: Buchmann et al. [6] employ purely hand-crafted solutions, which do not guarantee any correctness. Especially, they do not show that the synchronization result is still in the given modeling language. Some approaches such as [6, 9, 25, 31, 34, 37–39] do not consider conflicts between changes on both sides and/or do not provide the means to identify and solve them, which can lead to problems when model changes are overwritten by propagation without asking the modeler. Some propagation-based approaches are able to tackle the problem of detecting conflicts between parallel updates such as [16, 17, 19, 22, 25, 35, 39] by analyzing if a propagation step contradicts with a model change. However, as shown by Orejas et al. [29], propagation-based conflict detection is not deterministic in general and may depend on the order in which changes are propagated. This means that, depending on the propagation order, certain conflicts may or may not be detected. Other approaches are limited to a confluent set of grammar rules [19] or are limited to specific kinds of models such as tree-like hierarchies [31].

Constraint-based approaches are often based on a relational specification which can be enforced using tools such as a SAT solver. They typically solve problems globally; this means that all possible synchronization solutions are encoded in a search space. The approach proposed by Macedo et al. [28], for example, finds the closest model that is consistent again. Closeness can either be defined w.r.t. graph edit distance or be based on user defined distance metrics to support user-preferences. However, this flexibility comes at the price of scalability as constraint-based approaches can often cope with rather small models only.

Search-based approaches explicitly explore and find a rich set of synchronization solutions, which can become very expensive with increasing search space. Cicchetti et al. [8] align their work with that by Macedo et al. in that they calculate all closest sub-models that are still consistent to a given relational consistency specification. Focussing on sub-models, there is a potentially large amount of information loss as their approach does not truly incorporate all kinds

221 of model changes. Orejas et al. [30] propose a TGG-based
 222 approach, where a set of consistency-describing grammar
 223 rules is used to find all possible parse trees of the given
 224 inter-related models and enriching them with annotations
 225 for, e.g., mandatory, removed, added, and no longer covered
 226 elements. These annotations are used to find conflicts, which
 227 are resolved using back-tracking to calculate all possible syn-
 228 chronization solutions and to present them to the modeler to
 229 choose from. However, finding all possible synchronization
 230 solutions is very expensive and the amount of presented
 231 alternative solutions to the user might be overwhelming.

232 In summary, all the approaches discussed have one or
 233 more limitations. We are looking for a concurrent synchro-
 234 nization approach that (1) does not come with severe restrictions
 235 concerning the structure of the processed models or
 236 the definitions of the regarded consistency relation, (2) finds
 237 all kinds of conflicts between concurrent model changes
 238 in a deterministic way, (3) allows the modelers to interact
 239 with the synchronization process, (4) reliably returns a syn-
 240 chronization result that belongs to the given modeling lan-
 241 guage, and (5) scales along the size of model changes and
 242 conflicts rather than along the model size. In this paper, we
 243 will present an approach that has all these properties. How-
 244 ever, its scalability comes with the price of the restriction
 245 that consistency restoring operations modify only model
 246 parts that are causally dependent on those parts that are
 247 directly changed by modelers.

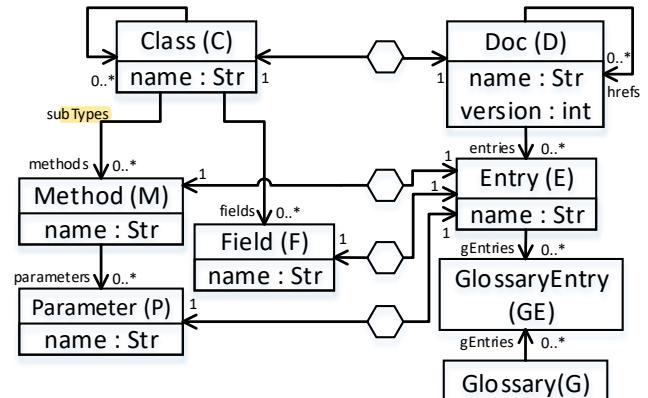
249 3 Triple Graph Grammars

250 In this section, we recall *triple graph grammars* (TGGs) [32],
 251 a declarative and rule-based approach to specify the con-
 252 sistency between two modeling languages. Being based on
 253 (typed attributed) graphs and their transformations as un-
 254 derlying formalism, TGGs are expressive and allow for the
 255 development of synchronization solutions with strong for-
 256 mal guarantees [11, 14, 21, 30]. Moreover, (many of) the
 257 operations needed during model synchronization algorithms
 258 can automatically be derived from the rules of a given TGG.
 259 Still, they can be implemented in a scalable way [1, 2]. As
 260 (typed attributed) graphs provide a suitable basis to formalize
 261 models and their transformations [3], TGG-based synchro-
 262 nization approaches are directly applicable to models. We
 263 thus use the terms “graph” and “model” interchangeably. In
 264 the following, the TGG concepts are recalled informally; a
 265 formal introduction can be found in, e.g., [10, 11]. An infor-
 266 mal introduction to graph transformation [18] (including a
 267 chapter on model translation and synchronization) appeared
 268 recently. We illustrate TGGs and the basic ingredients for
 269 our synchronization approach using a running example.

270 **Triple graphs and rules.** A *triple graph* consists of three
 271 graphs: a *source graph*, a *target graph*, and a *correspondence*
 272 *graph* in between that connects source and target graphs
 273 via two graph homomorphisms. The correspondence graph
 274

275 serves to establish traceability links between correlated el-
 276 ements from source and target graphs. In practical applica-
 277 tions, the underlying graphs are usually typed and attribut-
 278 ed. During synchronization processes, the occurring ob-
 279 jects may become *partial triple graphs* [13, 26]: A user may
 280 have deleted an element that was referenced by a correspon-
 281 dence morphism. Partial triple graphs still consist of three
 282 graphs; the graph homomorphisms connecting the corre-
 283 spondence graph with the source and target graph, however,
 284 may be partial, i.e., containing dangling references..

285 As a running example for a TGG, we define the consistency
 286 between a Java abstract syntax tree (AST) model (source)
 287 and a documentation model (target) as depicted in Fig. 2.
 288 This figure shows a metamodel (represented as a triple type
 289 graph) that declares the general syntax of models. The Java
 290 AST model consists of (Sub-)Classes containing Methods
 291 with Parameters and Fields, while the documentation model
 292 consists of Doc(ument)s with hyper references (href) to
 293 other Docs. Furthermore, a Doc contains Entries referencing
 294 Glossary Entries, that again are contained in a Glossary.
 295 Note that some elements have a name attribute, while Docs
 296 additionally store their version number. The correspondence
 297 types are depicted as hexagons referencing types of both,
 298 the Java and the documentation model, pair-wisely.

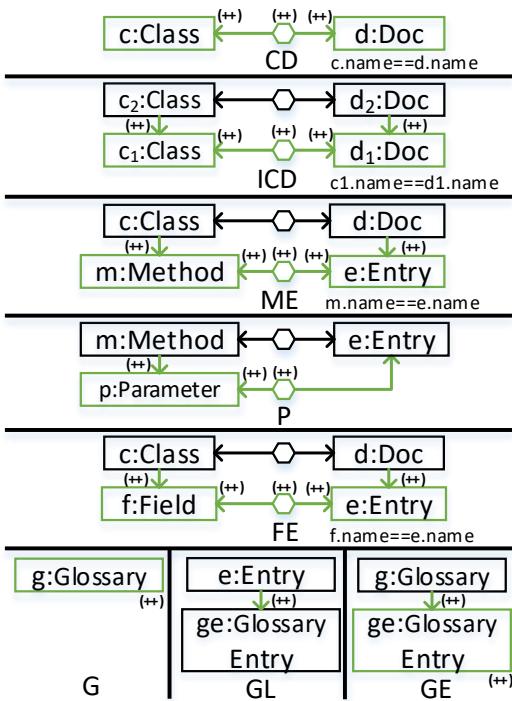


315 **Figure 2.** Metamodel of the running example

316 A *triple rule* consists of two triple graphs L and R (typed
 317 over the given triple type graph), called *left-hand side* (LHS)
 318 and *right-hand side* (RHS), respectively, such that their in-
 319 tersection $K = L \cap R$ is a triple graph again. Intuitively,
 320 the difference between L and K specifies all the elements to
 321 be deleted by an application of the rule, and the difference
 322 between R and K specifies all the elements to be created. Ad-
 323 ditionally, rules may be equipped with *negative application*
 324 *conditions* (NACs) that specify forbidden context in whose
 325 presence the rule is not applicable. Finally, we allow rules to
 326 be equipped with constraints concerning the attribute val-
 327 ues. In this paper, we restrict them to be equations involving
 328 attribute values of corresponding elements only.

331 A *triple graph grammar* consists of a start graph (that is
 332 usually the empty graph) and a set of non-deleting rules, i.e.,
 333 rules, where the LHS is a sub-triple graph of the RHS. The
 334 *language* defined by a TGG consists of all graphs derivable by
 335 an application sequence using its rules beginning at its start
 336 graph. Hence, a pair of models is consistent w.r.t. a given
 337 TGG iff a correspondence graph exists that extends the two
 338 models to a graph triple in the TGG language. Rewriting
 339 of partial triple graphs can be introduced analogously [26]
 340 and will be used to capture the semantics of synchronization
 341 operations manipulating dangling references.

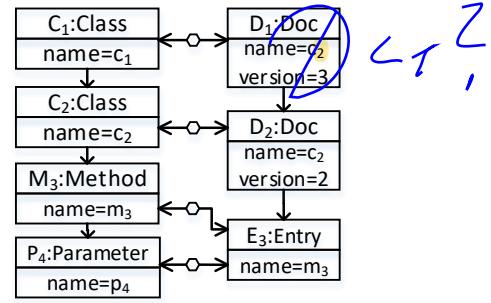
342 Figure 3 depicts the rule set of our running example con-
 343 sisting of 8 TGG rules. They are displayed in an integrated
 344 fashion, i.e., as a single graph. The black, unmarked elements
 345 constitute the LHS of the rule, i.e., the context that has to
 346 exist for a rule to be applicable. Green elements annotated
 347



371 **Figure 3.** TGG Rules

372 with (++) are to be created when the rule is applied. The rule
 373 *CD* has no precondition and thus may be used to generate
 374 Classes with corresponding Docs arbitrarily often. The rule
 375 *ICD* creates a Sub-Class and a corresponding Doc with a
 376 hyper-reference when a Class with a corresponding Doc
 377 already exist. Note that we create the subClass link together
 378 with the sub-class, which implicitly forbids multiple inher-
 379 itance. The rules *ME* and *FE* create Methods, resp. Fields,
 380 with corresponding Entries. Rule *P* creates a Parameter that
 381 corresponds to an already existing Entry. Finally, the rules *G*,
 382 *GE*, and *GL* create a Glossary together with Glossary Entries

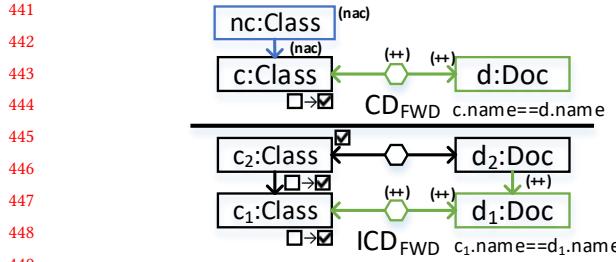
383 and links from Entries to Glossary Entries. These rules only
 384 act on the documentation model (the target side); the created
 385 elements do not have corresponding Java elements. Several
 386 rules (*CD*, *ICD*, *ME*, and *FE*) are equipped with attribute con-
 387 ditions. In each case, the condition declares that the names of
 388 the newly created elements should be equal. Figure 4 depicts
 389 a simple model that can be created applying first rule *CD*
 390 followed by applications of rules *ICD*, *ME*, and *P*.



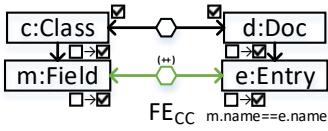
406 **Figure 4.** Exemplary Model

407 **Rules derived from TGG rules.** Triple rules can be used
 408 to create consistent models from scratch. For scenarios like
 409 *model translation* and *model synchronization*, suitable kinds
 410 of rules can be derived from the given TGG. First, a TGG rule
 411 can be *operationalized* to support forward (source → target)
 412 and backward (target → source) translations. Figure 5 de-
 413 picts the forward operationalized rules *CD_{FWD}* and *ICD_{FWD}*,
 414 which are created by converting all green source elements to
 415 context first as we consider them to already exist. To prevent
 416 elements from being translated twice, we introduce anno-
 417 tations: $\square \rightarrow \checkmark$ indicates that this element is still untranslated
 418 and applying the rule marks this element as translated. Con-
 419 sequently, \checkmark indicates that the annotated element must be
 420 translated before applying this rule. Note that *CD_{FWD}* con-
 421 tains a NAC, depicted in blue and annotated with (nac), that
 422 forbids the Class *c* to be translated into a corresponding Doc
 423 when it is a Sub-Class of another Class. This kind of NAC is
 424 also called a *filter NAC* [20, 24] and adjusts the translation
 425 process to avoid dead-ends. If a Sub-Class is translated using
 426 *CD_{FWD}* (without that NAC), for example, there is not any
 427 other forward rule that translates the remaining link to its
 428 Super-Class. The creation of all other forward and backward
 429 rules can be done analogously.

430 Another useful operationalization of TGG rules is referred
 431 to as *consistency check rules*. They detect whether yet un-
 432 translated elements in a source and a target model can be
 433 considered as correlated. If so, they create correspondence
 434 links. While these kinds of rules have been used to compute
 435 (maximal) correspondence relations between previously un-
 436 related models from the source and target domain [27], we
 437 employ them only *locally* to detect whether independently

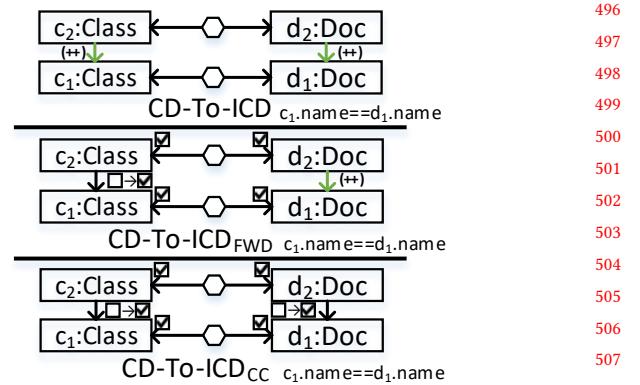
**Figure 5.** Exemplary TGG Forward Rules

453 added elements on source and target side can be considered
454 as corresponding to each other. We refer to this process as
455 *local-CC*. Figure 6 depicts an exemplary consistency check
456 rule that is derived from rule *FE*. Moreover, we project these
457 consistency check rules to their source and target parts only
458 to obtain *source* and *target patterns* which we will use for
459 conflict detection purposes later on.

**Figure 6.** Exemplary TGG Consistency Check Rule

469 Finally, there are *short-cut rules* that are synthesized from
470 TGG rules by means of a special kind of sequential rule com-
471 position operator [12]. Applying a short-cut rule replaces
472 one TGG rule application by another one while allowing to
473 preserve selected elements (instead of deleting and re-
474 creating them). Short-cut rules allow for advanced editing of
475 models while preserving information in the process. For-
476 ward or backward operationalizing a short-cut rule results
477 in *repair rules* that allow to directly propagate the edit the
478 short-cut rule specifies on the source side to the target (and
479 conversely). The crucial point is that short-cut rules spec-
480ify complex (language-preserving) edits one cannot imme-
481diately perform using the original rules of the given TGG.
482 Hence, their derived repair rules are often suitable to directly
483 propagate “free” user edits. Short-cut rules do not need to
484 be non-deleting and their derived repair rules may act on
485 partial triple graphs. For details on their construction, con-
486 ditions on language-preserving applications, and application
487 in unidirectional model synchronization, we refer to [12, 13].

488 An example of a short-cut rule is given in Fig. 7, where the
489 short-cut rule *CD-To-ICD* transforms an application of *CD* to
490 one of *ICD*. *CD-To-ICD_{FWD}* shows the forward operational-
491 ized short-cut rule, which directly propagates an edge that
492 was newly inserted between two classes. Furthermore, also
493 short-cut rules can be operationalized to obtain consistency
494 check operations such as *CD-To-ICD_{CC}*.

**Figure 7.** (Operationalized) Short-Cut Rule

Precedence graph. Given a consistent triple as depicted in Fig. 4, we can infer a *precedence graph* (PG) that describes with which TGG rule applications this triple can be derived and how those depend on one another. This means that a PG describes the aforementioned causal dependency relationship for model elements. Formally, a precedence graph is based on so-called *consistency patterns*. A consistency pattern is just the RHS of one of the rules of the given TGG. Applying a rule to a triple graph, one obtains a homomorphism from the RHS of the rule to the resulting triple graph: This homomorphism maps the elements of the LHS of the rule to the elements they have been matched to for applying the rule. The elements of its RHS that do not belong to the LHS (i.e., the elements to be created) are mapped to the newly created elements. Thus, given a sequence of rule applications, one obtains a family of homomorphisms from the rules’ RHSs to the triple graph. These homomorphisms are in a natural *dependency relation*: One homomorphism is dependent on another one, if the former matches an element, the underlying rule application of the latter one creates. Moreover, these homomorphisms *cover* the triple graph in the following sense: Every element of it is matched exactly once by a rule element to be created, i.e., by an element of RHS \ LHS of one of the rules. We will also say that this match *explains* or *accounts for* the element. Conversely, given a TGG and some triple graph, if there is a family of homomorphism from the consistency patterns to this triple graph such that the dependency relation is acyclic and the family *covers* the triple graph in the above sense, the triple graph belongs to the language of the TGG (for a formalization and proof, we refer to [27, Lemma 4]). This is what we define to be a *precedence graph* (PG) for a triple graph *M*: an acyclic graph where the nodes are homomorphisms from consistency patterns to *M* such that *M* is covered by them, and edges are their dependencies. Note that the dependency relation stored in a PG induces a causal dependency relation on elements of the triple graph: An element *x* (node or edge) depends on an element *y* if the element *y* is matched by the rule that

creates x . Similar information (about dependency between and coverage of elements) has been used by Kehrer to lift atomic model changes to the level of *edit scripts* [23].

Figure 8 depicts the PG for the model in Fig. 4 with each node corresponding to a TGG rule application where the name is based on the TGG rule name and an index representing the indices of created elements. The boxes inside the nodes represent the state of created elements of the corresponding rule application on source (left box) and on the target (right box) side. In the next chapter, we will introduce annotations for these boxes that will help us detecting conflicts.

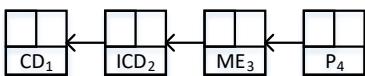


Figure 8. Exemplary Precedence Graph

4 Conflict detection

In this section, we present our approach to the detection of conflicts during model synchronization. We assume the following general setting (compare Fig. 1 again): A TGG is fixed; it defines a consistency relation between two modeling languages. Our concurrent synchronization process starts with a pair of consistent models M_1, M_2 , or, somewhat more formally, a triple graph belonging to the language of the given TGG (i.e., M_1 is the source and M_2 the target graph). This pair of consistent models comes with a precedence graph PG .

Both models are changed independently by two modelers, resulting in models M'_1 (source graph) and M'_2 (target side). Compared to M_1 , in M'_1 some elements may have been added, some deleted, and some attribute values may have changed; the same holds for M_2 and M'_2 . We call this change a *(model) delta*, and speak of *source* or *target delta* when we want to refer to only one of these changes. We do not make any assumptions on how they have been performed; we only assume that there is a way to identify the remaining elements of M_1 and M_2 with their counterparts in M'_1 and M'_2 , respectively.

Our goal is to find models M''_1 and M''_2 that are consistent, i.e., which are source and target graphs of a triple graph of the given language. Moreover, M''_1 and M''_2 should not differ too much from M'_1 and M'_2 , respectively. In general, there is no unique solution for this problem, even if requiring the distance between (M''_1, M''_2) and (M'_1, M'_2) to be *minimal* (according to some metric). Therefore, we provide modelers with the possibility to orchestrate the synchronization process, leading to individually defined outcomes.

To compute a pair of consistent models, we first extend and annotate the precedence graph to obtain a *delta precedence graph* (DPG). Basically, this graph comprises information

which parts of the PG have been affected by the deltas and how the individual changes can (locally) be propagated to the respective other side. *Our guiding principles for synchronization are to only change the models where the delta makes this necessary and to preserve as many of the model changes as possible, e.g., to not delete newly added elements and to not recreate deleted ones* (compare, e.g., [7, 33] for these principles). Whenever it is not possible to simultaneously preserve user-changes on the source and on the target model such that the model remains in the language of the TGG, we call this a *conflict*. Our synchronization process first analyzes the delta precedence graph for conflicts and subsequently, propagates the source and target deltas to the respective other side while resolving the detected conflicts according to an orchestration given by the user.

In this section, we introduce delta precedence graphs and different types of conflicts illustrated using our running example.

Figure 9 shows an example of a concurrent model change applied to a formerly consistent model graph. In the source graph of the original model, there are two Classes (C_1 and C_2) that have a Field and Method each. Furthermore, M_6 has two Parameters and M_8 has one. The target graph has two Docs (D_1 and D_2) that contain two Entries each, where the Entries contained in D_2 reference the Glossary Entry GE_{13} . GE_{13} and GE_{12} are contained in the Glossary G_{14} . Elements with the same indices on both sides have a correspondence link, except for Parameters, which are connected to those Entries that correspond to their Method.

Several concurrent *model changes* have taken place with different impacts and issues. First, the Field F_7 is pulled up into a newly created Super-class C_3 of C_2 . In the target graph, the corresponding Entry E_7 is moved to D_1 , which does not correspond to C_3 . Also, a new Field F_4 and a new Entry E_4 are created within elements C_1 and D_1 , respectively. They have the same names. Additionally, M_6 is deleted together with its Parameter P_9 , while the Parameter P_{10} was moved to M_8 . However, while M_6 is deleted in the source graph, the corresponding target element E_6 as well as E_4 , E_5 , and E_7 are linked to GE_{12} . Finally, the names of corresponding model elements M_8 and E_8 are changed to different values. Similarly to our notation of rules, in Fig. 9, newly added elements are depicted in green and marked with $(++)$, deleted ones in red and marked with $(--)$, and attribute changes are indicated via an arrow \rightarrow .

4.1 Delta Precedence Graphs

A *delta precedence graph* extends a precedence graph by information on how a delta affected the validity of the precedence graph. For elements added in the delta, we need to find suitable TGG-rule applications that could have created these elements; potentially correlating elements have to be created on the other side. Moreover, consistency matches of

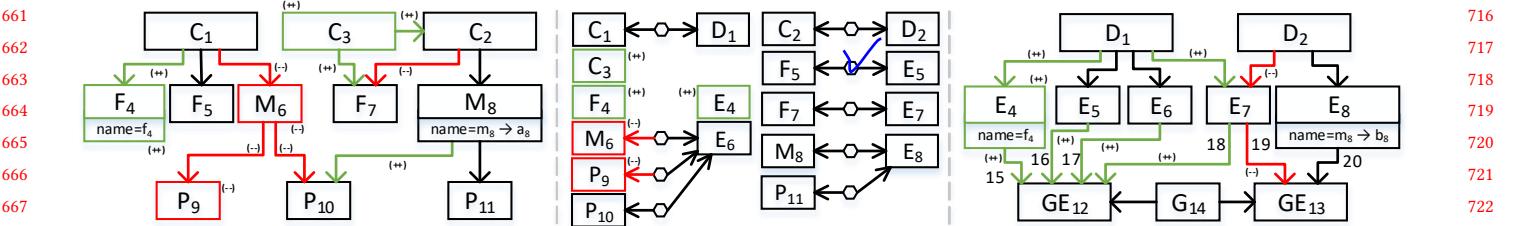


Figure 9. Running Example – Model and Delta

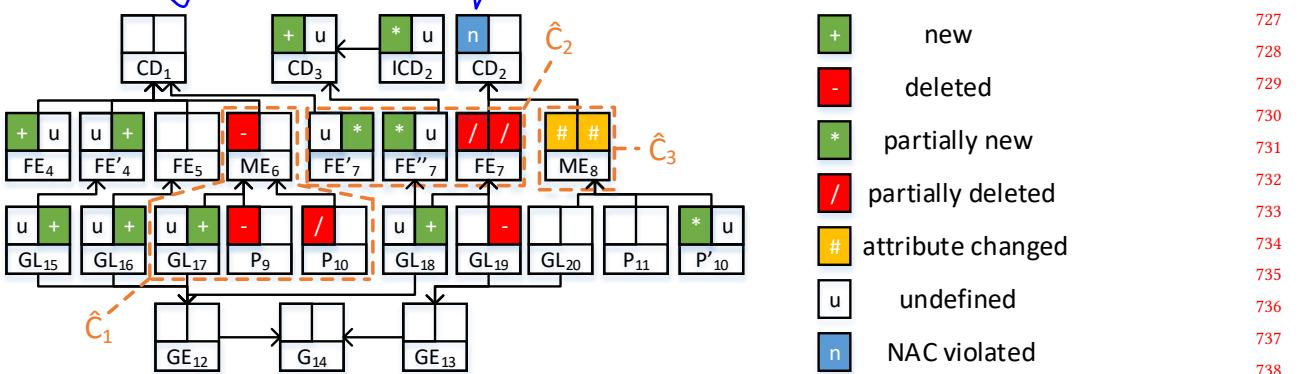


Figure 10. Running Example – Delta Precedence Graph

the precedence graph may have been invalidated (or “broken”) in three different ways: (i) attribute values have been changed such that an attribute condition of a TGG rule is violated, (ii) elements that are covered by a consistency match have been deleted, and (iii) elements have been added so that already existing elements have to be parsed anew since a NAC is violated. Elements that are no longer covered by a consistency pattern since the formerly covering one is “broken” have to be matched anew as well. All this information is collected in a DPG. In the definition of DPG, we call an element of the updated model *unpropagated* if one of the following two cases applies:

1. This element has been newly added by one of the user edits, or
2. there is a consistency pattern that explains how this element has been created (i.e., the underlying rule application created this element). However, this rule application has been rendered invalid by the user edits (either, because one of its filter NACs is violated now or because a matched element has been deleted).

Definition 4.1 (Delta precedence graph). Let a TGG and a pair of consistent models (M_1, M_2) together with a precedence graph PG for it be given. The *delta precedence graph* (DPG) for (M_1, M_2) , PG , and a delta (consisting of sequences of graph changes leading from M_1 to M'_1 and M_2 to M'_2 , respectively) consists of the nodes of the given precedence

graph PG and a set of new nodes consisting of matches for source and target patterns that are such that at least one unpropagated element is matched by an element the underlying rule of the pattern creates.

Edges are, again, defined via dependencies.

Moreover, the nodes of the DPG have *source* and *target* annotations over the alphabet $\{+, -, *, /, \#, u, n\}$ according to the following rules:

- Nodes stemming from the original PG, i.e., nodes indicating formerly valid rule applications, are annotated on the source side in the following way:
 - annotation “-” whenever all elements the underlying rule application created on the source side have been deleted by the source edit;
 - annotation “/” whenever some (but not all) elements the underlying rule application created on the source side have been deleted by the source edit or if the match for this consistency pattern is broken because context elements are missing;
 - annotation “#” whenever some attribute values were changed by the source edit such that at least one attribute constraint of the underlying rule application is violated;
 - annotation “n” whenever the source edit added a new element that introduces a violation of a source NAC of the forward rule of the rule from which the consistency pattern is derived.

The target annotation is defined completely analogously. A node may also have more than one such annotation, which is why we will refer to sets of annotations in the following. However, for simplicity reasons, nodes in our running example only have sets of size 1. ✓

- Whenever all nodes that the underlying rule created on the source side are matched to unpropagated elements, a node corresponding to a match for a source pattern is annotated with “+” on the source and with “u” on the target side. Symmetrically, the annotations for such a target pattern are switched.
- Whenever at least one node on the source side, created by the underlying rule, is matched to an already propagated element that was part of another formerly intact consistency match, a node corresponding to a match for a source pattern is annotated via “*” on the source and “u” on the target side. Symmetrically, the annotations for such a target pattern are switched.

Figure 10 depicts the DPG that corresponds to the model changes described above. One of the nodes FE_7 , for example, is marked with “/” on both source and target side because the edges created by the corresponding rule application have been deleted on both sides; however, nodes F_7 and E_7 are preserved. The two new nodes indexed with FE'_7 and FE''_7 denoting a new source and target pattern match, respectively, indicate by which rule the now unaccounted nodes F_7 and E_7 could possibly have been created. The consistency restoration has to check whether they can be combined into a single TGG-rule application creating the two correlated elements simultaneously. Node ME_8 is annotated with “#” on both sides as the attribute values have been changed on both sides and the constraint requiring equal names is violated now. As a last example, the node CD_2 is annotated with “n” on the source side: Due to the newly introduced inheritance edge, it is no longer possible for the class C_2 to be created using rule CD . Thus, this node becomes unpropagated; the node indexed with ICD_2 indicates a new possibility to parse this node. ✓

4.2 Conflicts

As discussed above, the annotations of a DPG indicate some synchronization actions that have to take place; as long as the dependency is respected (i.e., is still acyclic), a consistent triple graph is restored as soon as every annotation has been dealt with. Our definition and treatment of conflicts is based on the already mentioned idea of change-preservation. Whenever a user edit directly affected an element, this effect should be preserved. This means that newly created elements or elements whose attribute values have been changed are *intended to persist*. Deleted elements are intended to remain deleted. A conflict is a situation where all options available to propagate a certain change require to undo another one. We ✓

classify such conflicts based on the annotations occurring in DPGs and use the term *conflict scope* to refer not only to a conflict but also to other elements that depend on how the conflict is resolved. In the following, we assume a pair of models that was originally consistent, a PG for it, a model delta, and the induced DPG to be given. We illustrate all kinds of conflicts using our example in Fig. 10 (where the *conflict scopes* \hat{C}_1 , \hat{C}_2 , and \hat{C}_3 are indicated by dashed orange lines). ✓

Preserve-delete conflict. A preserve-delete conflict is a situation where one of the deltas deletes a certain element whereas its corresponding element is used in the other delta and thus, intended to persist: A *potential preserve-delete conflict* is a node of the DPG where “-” or “/” belongs to the source or target annotation (except for the case where both annotations are “-”). It is a *preserve-delete conflict* if there is a newly added element on the other side or an element whose attribute value has been changed such that, by propagating the deletion, all patterns that would have been able to create that element vanish. Its *scope* is the node itself and all nodes transitively depending on it. ✓

\hat{C}_1 depicts a *preserve-delete conflict* that is characterized by a full deletion (“-”) in one precedence node on the source side and a creation of new elements (“+”) in one of its dependent nodes on the target side. ✓

Here, M_6 is deleted on the source side, but a new reference is added between E_6 and GE_{12} on target side. The conflict scope \hat{C}_1 includes changes that must not directly be in conflict with each other, but have to be considered when resolving the conflict, e.g., non-revoking P_9 implies non-revoking M_6 . ✓

Correspondence preservation conflict. A correspondence preservation conflict is a situation, where deltas modify corresponding elements on both source and target side such that it is not possible to restore the consistency without either deleting the correspondence relationships between the affected elements (and creating new correspondence relationships to different elements, which in the general case have to be created, too) or without discarding changes in the source or target model. In our example, \hat{C}_2 is such a conflict where F_7 is moved to C_3 while E_7 is moved to D_1 . However, C_3 and D_1 do not correspond to each other and thus there is no TGG consistency pattern that can correlate these changes. Hence, we have to decide whether to revoke the relocation of either F_7 or E_7 . ✓

Every node of the DPG where “n” or “/” belongs to the source and target annotation is a *potential correspondence preservation conflict*. It is a *correspondence preservation conflict* if there exists no intact TGG consistency pattern match that again covers and relates these elements. Its *scope* is the node itself as well as nodes with a “*” annotation whose ✓

881 corresponding rule applications create some elements that
 882 are also created by the rule application of the node itself.

883 **Attribute change conflict.** An attribute change conflict
 884 is a special case of a correspondence preservation conflict,
 885 where attribute values of until now corresponding source
 886 and target elements have been changed in such a way that
 887 the attribute values on both sides are no longer consistent. In
 888 our example, \hat{C}_3 is an attribute change conflict that occurs in
 889 node ME₈. The name of M₈ is changed from m₈ to a₈, while
 890 that of E₈ is changed from m₈ to b₈.

891 If both names would have been changed equally, no con-
 892 flict would have been detected and the incremental pattern
 893 matcher would not detect any broken rule application. Ev-
 894 ery node of the DPG where “#” belongs to the source and
 895 the target annotation is a *potential attribute change conflict*.
 896 It is an *attribute change conflict* if furthermore an attribute
 897 constraint is violated because both attribute values were
 898 changed.

899 In Appendix A, we provide a overview of how the annota-
 900 tions of a node in a DPG relate to possible conflicts.

5 Consistency restoration

901 In this section, we will present a catalog of concurrent syn-
 902 chronization fragments, which can be orchestrated and ex-
 903 ecuted sequentially to restore the consistency of the model
 904 under change in a concurrent synchronization scenario. This
 905 catalog includes the previously introduced operationaliza-
 906 tions of TGG rules and three pre-defined conflict resolution
 907 strategies. Users can orchestrate their specific consistency
 908 restoration processes using these fragments.

909 In the following, we present all the concurrent synchro-
 910 nization fragments of our catalog, show an example orches-
 911 tration of them which is quite typical, and apply this orches-
 912 tration to our running example. All the following fragments
 913 except for *Resolve Conflict* are applied to elements only that
 914 do not belong to a conflict. *Resolve Conflict* is applied to each
 915 conflict (scope). Furthermore, if a fragment is applied, it pro-
 916 cesses all feasible elements until no unprocessed one can be
 917 found.

- 918 • **Local CC** is used to find and correlate newly added
 919 elements in the source and target graphs that may
 920 correspond to each other. This fragment processes
 921 related pairs of precedence nodes annotated with (+|u)
 922 and (u|+). If it is chosen, it has to be applied before
 923 *Translate* as this also processes newly added elements
 924 that are no longer available afterwards.

- 925 • **Translate** is used to translate newly added elements
 926 to the opposite side and thus complete the rule appli-
 927 cation by applying a forward or backward operational-
 928 ized TGG rule. This fragment processes precedence
 929 nodes annotated with (+|u) and (u|+).

- 930 • **Repair** employs short-cut rules to fix broken prece-
 931 dence nodes. Forward and backward operationalized

932 short-cut rules propagate complex changes from one
 933 side to the other. Consistency check operationalized
 934 short-cut rules allow to find corresponding changes
 935 on both sides and resolve them if possible. This frag-
 936 ment processes precedence nodes that are annotated
 937 with “n” or “/” on one or both sides and related nodes
 938 annotated with (*|u) or (u|*).

- 939 • **Resolve conflict** is applied to each conflict (scope)
 940 and can be configured for each type of conflict that we
 941 identified in the previous chapter. We can call *Trans-*
942 late, *Repair* and *Propagate*, where *Repair* can be used
943 to reduce the conflict size beforehand, while *Trans-*
944 late and *Propagate* can only be used after the conflict
945 (scope) has been resolved because both include propa-
946 gation steps that can only be executed afterwards. For
947 resolving a conflict (scope), we offer three pre-defined
948 strategies:

- 949 – **Take Source** discards all the changes on the target
 950 side.
- 951 – **Take Target** discards all the changes on the source
 952 side.
- 953 – **Preserve** discards all the deletions that block newly
 954 added elements from being propagated.

955 All applied fragments within *Resolve Conflict* are ap-
 956 plied to elements only that belong to the current con-
 957 flict (scope).

- 958 • **Rollback** revokes rule applications in cases where all
 959 the deletions were performed consistently on both
 960 sides or only on one side and the other side remains
 961 untouched, e.g., all green source elements were deleted.
 962 These nodes are annotated solely with “-” on one or
 963 both sides.

- 964 • **Propagate** applies *Repair* first to fix broken matches
 965 rather than revoking them. Then, it applies *Rollback* to
 966 revoke rule applications that have been consistently
 967 deleted on one side. Finally, *Translate* is applied and
 968 translates newly added elements to the opposite side.
 969 From our experience, calling these fragments in that
 970 specific order is a good choice for another fragment
 971 as it yields a sequential synchronization control flow.

- 972 • **Clean up** deletes all elements from source, correspon-
 973 dence and target graphs that are currently inconsistent
 974 w.r.t. our TGG. This fragment can only be called at the
 975 end of a synchronization process but can also be omit-
 976 ted if the user decides to eliminate inconsistencies in
 977 this way later on.

978 To restore the consistency of the model in our running exam-
 979 ple, we choose the following orchestration of fragments to
 980 be applied in the given sequential order: *Local CC* → *Trans-*
981 late → *Repair* → *Resolve Conflict* {*Repair* → *Take Source*
982 → *Propagate*} → *Propagate* → *Clean up*. For simplicity rea-
983 sons, we resolve all three conflicts (\hat{C}_1 , \hat{C}_2 , \hat{C}_3 from Fig. 10)
984 uniformly.

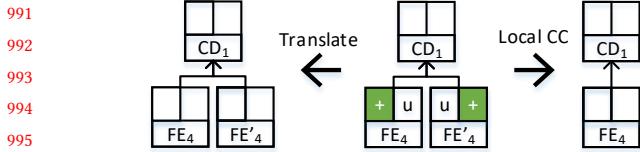


Figure 11. Running Example – Translate vs. Local CC

To correlate corresponding changes on both sides, we apply *Local CC* first. Nodes F_4 and E_4 are created independently with the same name but not set into correspondence yet. Using *Local CC*, we are also able to relate both changes to each other, creating the missing correspondence link in between. The effect on the precedence graph is shown on the right of Fig. 11. On the left of this figure, we see the result of not using *Local CC* but *Translate*. Applying *Translate* on both F_4 and E_4 independently of each other would create corresponding elements on the opposite sides using the proper forward and backward rule. This step results in new elements F'_4 and E'_4 besides the former ones.

Next, we apply *Translate* to several changes that are not contained in a conflict (scope) (as described in Section 4). This is the case for CD_3 , GL_{14} , and GL_{15} . Translating newly added elements establishes consistency of each precedence node since the underlying rule applications are now complete and thus consistent.

Afterwards, we apply *Repair* to fix broken rule applications such as CD_2 that is inconsistent due to a NAC violation. Figure 7 depicts the shortcut rule $CD\text{-To-}ICD_{FWD}$ that can be applied here to transform CD_2 into ICD_2 . The effect is that D_2 is preserved and an edge between D_3 and D_2 is created. The corresponding effect on the precedence graph is shown in Fig. 12. Note that the prior propagation of CD_3 is necessary to create the context needed by $CD\text{-To-}ICD_{FWD}$ to be applicable.

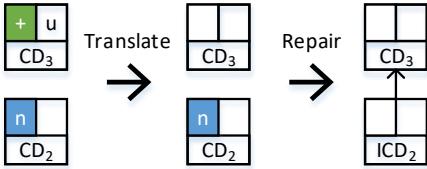


Figure 12. Running Example – Translation then Repair

Now, only conflicts (scopes) are still there, which are resolved in any order using *Resolve Conflict*. The intermediate model and delta precedence graph can be found in Appendix C. Starting with conflict scope \hat{C}_1 , the primary issue is that a glossary link was created at entry E_6 while deleting the method M_6 that corresponds to E_6 . Considering the precedence node P_{10} , only parts were deleted which implies that some of the remaining elements are to be preserved (here

P_{10}). Using *Repair* first allows us to reduce the size of \hat{C}_1 by propagating the re-location of P_{10} from M_6 to M_8 . This step is depicted in Fig. 13.

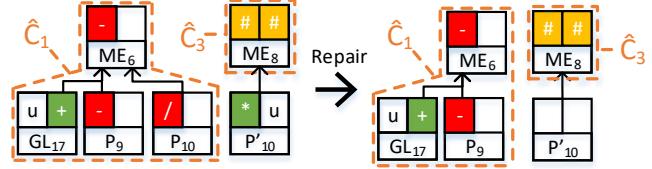


Figure 13. Running Example – Reduce conflict scope \hat{C}_1

Figures 14 (b) – (d) depict the results of all three conflict resolution strategies applied to conflict scope \hat{C}_1 in Figure 14 (a). The application of *Preserve* is of special interest as it revokes deletion deltas that block create deltas from being propagated. Following that strategy, the changes to ME_6 are revoked to solve the conflict, while keeping the changes to P_9 untouched. Applying any of these strategies leaves the remaining elements in a state where they can be propagated without colliding with any changes on the opposite side. To resolve \hat{C}_1 finally, we use *Take Source* (as specified earlier) and revoke all changes on the target side that are related to the conflict.

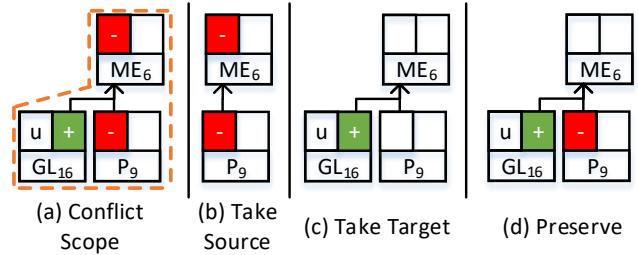
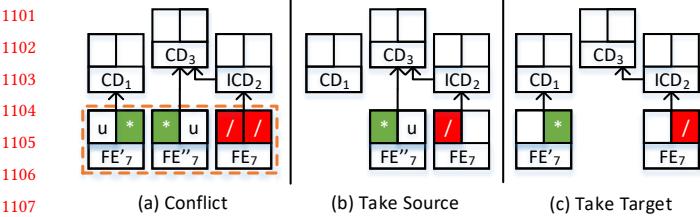


Figure 14. Running Example – Resolving \hat{C}_1

For resolving the conflict scope \hat{C}_2 , we perform *Repair* first, but this application does not have any effect here as both changes contradict each other. Figure 15 depicts the results of *Take source* (b) and *Take target* (c), which revoke the changes on one side. Using the operationalized shortcut rule $CD\text{-To-}ICD_{FWD}$ after (b) or $CD\text{-To-}ICD_{BWD}$ after (c) propagates the changes to the opposite side, respectively. Note that *Preserve* has no effect here due to the nature of *correspondence preservation* conflicts. Since we chose *Take Source* as conflict resolution strategy, we can react to the relocation of F_7 to C_3 now by also moving E_7 from D_2 to D_3 . This is done by applying *Propagate*.

Finally, we apply *Repair* to conflict scope \hat{C}_3 , which has no effect as both attribute changes contradict each other. Resolving \hat{C}_3 can be done by choosing either *Take source* or *Take target*. Again, *Preserve* would not have any effect since there

Figure 15. Running Example – \hat{C}_2

are no deletions that block the propagation of additions. Instead, we have to decide which attribute change to propagate (which implies to revoke the opposite one). Applying *Take Source* and subsequently *Propagate*, the remaining attribute change is propagated by re-evaluating the corresponding attribute constraint and transferring the value to E_8 .

This leaves *Clean up* with nothing to do since all changes have been accounted for. The final model and delta precedence graph are depicted in Appendix C.



6 Implementation

Our approach is implemented in a synchronization component as part of the state-of-the-art model transformation tool eMoflon [36]. Figure 16 depicts the synchronization components with its interdependencies to an incremental pattern matcher and its inputs and outputs. In our synchronization framework, we allow modelers to change the source and target of a triple graph independently. eMoflon keeps track of these changes by employing an incremental pattern matcher that throws events when new matches of TGG rules have been detected or existing ones have been invalidated. This information is used to update a (delta) precedence graph which represents the dependencies between TGG rule applications. The synchronization component analyzes the delta precedence graph to detect conflicts. They can be resolved by a synchronization and conflict resolution orchestration that has been implemented by an integration manager, who is an expert in both source and target domain. However, we offer pre-defined configurations such as the one from our running example in the previous chapter that can be extended. Applying this orchestration resolves the previously detected conflicts step-by-step and restores consistency of the triple graph. eMoflon already has an extensive test suite¹ with 344 tests for various TGG-based consistency restoration scenarios from which 25 constitute concurrent synchronization tests. In the near future, we will extend this test suite especially w.r.t. to more concurrent synchronization scenarios.

7 Evaluation

In this section, we will present our evaluation results, which are based on our implementation in eMoflon and two different TGG projects. The first TGG is our running example

¹<https://github.com/eMoflon/emoflon-ibex-tests>.

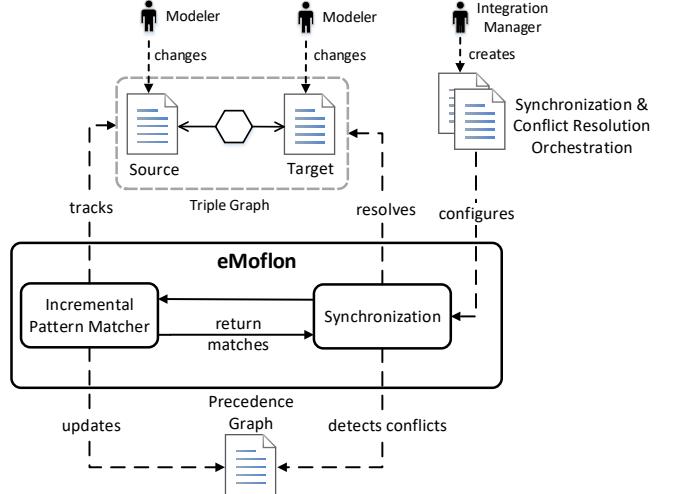


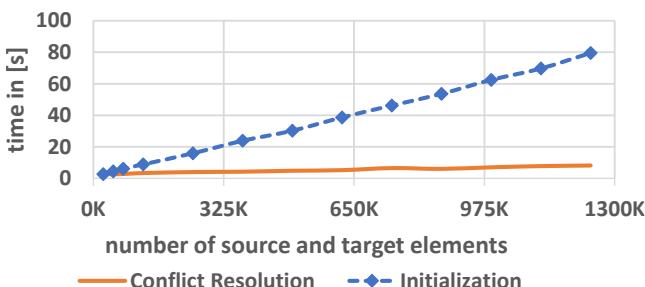
Figure 16. eMoflon - Synchronization Components

consisting of 8 TGG rules, while the second one is based on the example introduced in [13] and consists of 28 TGG rules that define consistency between MoDisco [5] and custom documentation models. However, the second TGG does not only contain more rules but also more asymmetric rules, i.e., rules that are no simple 1-to-1 mapping between source and target. As a test environment, we use a workstation with an AMD Ryzen Threadripper 2990WX 3GHz 32xCore using 128GB of RAM. One of the main goals of our approach is to provide a scalable concurrent synchronization solution. Therefore, we pose the following research questions: **RQ1:** Does our synchronization approach scale with the size of the model or with the number of changes and conflicts? **RQ2:** Does the performance of the conflict detection change if less changes lead to actual conflicts? **RQ3:** How does the number of changes and conflicts affect the conflict resolution performance?

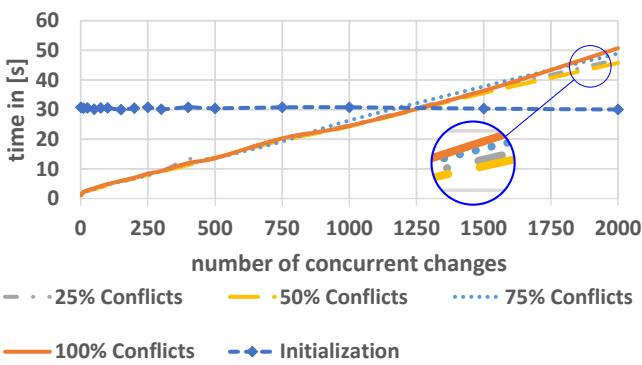
To answer these questions, we investigate two scenarios for each TGG: First, we generate a fixed number of 100 conflicts and increase the size of both, the source and the target model, which we measure in number of nodes. Second, we choose a fixed model size of about 500 000 nodes (in the source and target models together) and increase the number of concurrent changes. To investigate the correlation between changes and actual conflicts, we distinguish between 4 sub-scenarios by choosing the changes in such a way that 25%, 50%, 75% and 100% lead to an actual conflict, respectively. For all scenarios, we plot the initialization time where the incremental pattern matcher collects all matches for the still consistent triple graph, i.e., computes the precedence graph. Then, we apply a number of changes to both, the source and target model and measure the time to restore consistency. Note that each change is meant to induce one of the three conflict types previously presented. For each data point, we measured 20 repetitions and took the average

1211 value over all runs. Note that the plots of the second example
 1212 can be found in Appendix B.

1213 Figure 17 depicts the plot of the first scenario and our running
 1214 example: The initialization time increases for this TGG
 1215 linearly with the size of the model, while the conflict reso-
 1216 lution time stays almost constant.² The same holds for the
 1217 second TGG with a larger set of rules, which takes about 19 %
 1218 longer to initialize. This means that the model size does not
 1219 directly affect the performance of conflict resolution (**RQ1**).
 1220 Figure 18 depicts the plots for the second scenario and our
 1221 running example: The initialization time stays constant with
 1222 a constant model size of 500 000 (source and target) nodes
 1223 while the time to detect and resolve conflicts increases lin-
 1224 early. Whether the changes to both sides are in conflict with
 1225 each other has only a minor impact on the performance and
 1226 increases the gradient slightly. For the larger set of rules,
 1227 the impact becomes more significant and takes 16 % more
 1228 seconds per 25 % more conflicts (**RQ2**). Thus, we can con-
 1229 clude that the performance scales linearly with the size of
 1230 changes and to some point with the amount of conflicts that
 1231 are introduced by changes (**RQ3**).



1233 **Figure 17.** Increasing model size with constant number of
 1234 100 conflicts



1235 **Figure 18.** Increasing number of changes with constant
 1236 model of 500K nodes

1263 ²It slightly increases since deletions become more expensive with increasing
 1264 model size in the Eclipse Modeling Framework.

1266 **Threats to validity.** Our evaluation is based on synthe-
 1267 sized models and changes only. Thus, it remains future work
 1268 to investigate real-world scenarios by analyzing, e.g., Git
 1269 merge requests and deducing models from code. Further-
 1270 more, we evaluated only with two TGGs that have different
 1271 characteristics but describe a similar scenario. However, the
 1272 rules of both TGGs are symmetric as well as asymmetric and
 1273 thus, do not only represent simple 1-to-1 mappings, which
 1274 makes them representative for a broader range of TGGs.

8 Conclusion

1276 In this paper, we presented a scalable TGG-based, precede-
 1277 nce-driven concurrent model synchronization approach. We
 1278 showed how to use a delta precedence graph to identify con-
 1279 flicts and introduced a modular framework that enables de-
 1280 velopers to resolve these conflicts and, furthermore, orches-
 1281 trate the whole process to achieve specific synchronization
 1282 goals. Giving a broad overview of the landscape of concu-
 1283 rrent synchronization works, we showed that this approach
 1284 is indeed novel in that we detect conflicts before propagating
 1285 any changes. Furthermore, we showed how to detect a new
 1286 kind of conflict, namely *correspondence preservation* conflicts,
 1287 which to the best of our knowledge, no other approach is able
 1288 to detect so far. Our approach has been implemented and
 1289 evaluated within the state-of-the-art graph transformation
 1290 tool **eMoflon**. In the evaluation, we showed for two differ-
 1291 ent TGG projects that our synchronization approach scales
 1292 linearly with the size of changes instead of the model size.

1293 For the future, we plan to extend and formalize our ap-
 1294 proach w.r.t. handling further types of conflicts between rule
 1295 applications that, if applied, would violate, e.g., multiplicity
 1296 constraints of a metamodel. Further investigations are also
 1297 needed to study the effects of one major design decision of
 1298 our concurrent model synchronization approach in practice:
 1299 still consistent rule applications remain untouched if they
 1300 do not (causally) depend on a broken rule application. This
 1301 design decision is inspired by a *least change and least surprise*
 1302 principle [7] and is one main reason for the scalability of
 1303 our approach. Due to this, our synchronization algorithm is
 1304 in general not always able to reestablish the consistency of
 1305 two models when a change in one model requires its propa-
 1306 gation against the introduced causal dependency relation in
 1307 the related other model. However, there are works such as
 1308 [15, 24] that show how to derive application conditions with
 1309 a statically analyzable criterion for TGGs such that our syn-
 1310 chronization algorithm very rarely runs into such a situation.
 1311 In our experience TGGs violate this criterion if and only if
 1312 the propagation of a local change in one model would have
 1313 rather unexpected and unwanted global effects on the other
 1314 model from a user’s point of view. But further experiments
 1315 are needed to confirm our experiences. Finally, we plan to
 1316 build up a rich zoo of concurrent synchronization scenarios
 1317 and to compare different approaches with respect to these.

1321 References

- 1322 [1] Anthony Anjorin, Thomas Buchmann, Bernhard Westfechtel, Zinovy
1323 Diskin, Hsiang-Shang Ko, Romina Eramo, Georg Hinkel, Leila Samimi-
1324 Dehkordi, and Albert Zündorf. 2020. Benchmarking bidirectional
1325 transformations: theory, implementation, application, and assessment.
1326 *Journal of Software and Systems Modeling* 19, 3 (2020), 647–691. <https://doi.org/10.1007/s10270-019-00752-x>
- 1327 [2] Anthony Anjorin, Zinovy Diskin, Frédéric Jouault, Hsiang-Shang
1328 Ko, Erhan Leblebici, and Bernhard Westfechtel. 2017. BenchmarX
1329 Reloaded: A Practical Benchmark Framework for Bidirectional Trans-
1330 formations. In *Proceedings of the 6th International Workshop on Bi-
1331 directional Transformations co-located with The European Joint Confer-
1332 ences on Theory and Practice of Software, BX@ETAPS 2017, Uppsala, Swe-
1333 den, April 29, 2017 (CEUR Workshop Proceedings)*, Romina Eramo and
1334 Michael Johnson (Eds.), Vol. 1827. CEUR-WS.org, 15–30.
- 1335 [3] Enrico Biermann, Claudia Ermel, and Gabriele Taentzer. 2012. Formal
1336 foundation of consistent EMF model transformations by algebraic
1337 graph transformation. *Journal of Software and Systems Modeling* 11, 2
1338 (2012), 227–250. <https://doi.org/10.1007/s10270-011-0199-7>
- 1339 [4] Marco Brambilla, Jordi Cabot, and Manuel Wimmer. 2017. *Model-Driven Software Engineering in Practice, Second Edition*. Morgan & Claypool Publishers. <https://doi.org/10.2200/S00751ED2V01Y201701SWE004>
- 1340 [5] Hugo Brunelière, Jordi Cabot, Grégoire Dupé, and Frédéric Madiot.
1341 2014. MoDisco: A model driven reverse engineering framework. *Jour-
1342 nal of Information and Software Technology* 56, 8 (2014), 1012–1032.
1343 <https://doi.org/10.1016/j.infsof.2014.04.007>
- 1344 [6] Thomas Buchmann and Sandra Greiner. 2016. Handcrafting a Triple
1345 Graph Transformation System to Realize Round-trip Engineering
1346 Between UML Class Models and Java Source Code. In *Proceedings
1347 of the 11th International Joint Conference on Software Technologies
1348 (ICSOFT 2016) - Volume 2: ICSOFT-PT, Lisbon, Portugal, July 24 - 26,
1349 2016*, Leszek A. Maciaszek, Jorge S. Cardoso, André Ludwig, Marten
1350 van Sinderen, and Enrique Cabello (Eds.). SciTePress, 27–38. <https://doi.org/10.5220/0005957100270038>
- 1351 [7] James Cheney, Jeremy Gibbons, James McKinna, and Perdita Stevens.
1352 2017. On principles of Least Change and Least Surprise for bidirectional
1353 transformations. *Journal of Object Technology* 16, 1 (2017), 3:1–31.
1354 <https://doi.org/10.5381/jot.2017.16.1.a3>
- 1355 [8] Antonio Cicchetti, Davide Di Ruscio, Romina Eramo, and Alfonso
1356 Pierantonio. 2010. JTL: A Bidirectional and Change Propagating Trans-
1357 formation Language. In *Software Language Engineering - Third Inter-
1358 national Conference, SLE 2010, Eindhoven, The Netherlands, October
1359 12-13, 2010, Revised Selected Papers (Lecture Notes in Computer Science)*,
1360 Brian A. Malloy, Steffen Staab, and Mark van den Brand (Eds.), Vol. 6563.
1361 Springer, 183–202. https://doi.org/10.1007/978-3-642-19440-5_11
- 1362 [9] Alexander Egyed. 2007. Fixing Inconsistencies in UML Design Models.
1363 In *29th International Conference on Software Engineering (ICSE 2007),
1364 Minneapolis, MN, USA, May 20-26, 2007*. IEEE Computer Society, 292–
1365 301. <https://doi.org/10.1109/ICSE.2007.38>
- 1366 [10] Hartmut Ehrig, Karsten Ehrig, Ulrike Prange, and Gabriele Taentzer.
1367 2006. *Fundamentals of Algebraic Graph Transformation*. Springer.
1368 <https://doi.org/10.1007/3-540-31188-2>
- 1369 [11] Hartmut Ehrig, Claudia Ermel, Ulrike Golas, and Frank Hermann. 2015.
1370 *Graph and Model Transformation - General Framework and Applications*.
1371 Springer. <https://doi.org/10.1007/978-3-662-47980-3>
- 1372 [12] Lars Fritzsche, Jens Kosiol, Andy Schürr, and Gabriele Taentzer. 2018.
1373 Short-Cut Rules - Sequential Composition of Rules Avoiding Unneces-
1374 sary Deletions. In *Software Technologies: Applications and Foundations
1375 - STAF 2018 Collocated Workshops, Toulouse, France, June 25-29, 2018,
1376 Revised Selected Papers (Lecture Notes in Computer Science)*, Manuel
1377 Mazzara, Iulian Ober, and Gwen Salaün (Eds.), Vol. 11176. Springer,
1378 415–430. https://doi.org/10.1007/978-3-030-04771-9_30
- 1379 [13] Lars Fritzsche, Jens Kosiol, Andy Schürr, and Gabriele Taentzer. 2019.
1380 Efficient Model Synchronization by Automatically Constructed Repair
1381 Processes. In *Fundamental Approaches to Software Engineering - 22nd
1382 International Conference, FASE 2019, Held as Part of the European Joint
1383 Conferences on Theory and Practice of Software, ETAPS 2019, Prague,
1384 Czech Republic, April 6-11, 2019, Proceedings (Lecture Notes in Computer
1385 Science)*, Reiner Hähnle and Wil M. P. van der Aalst (Eds.), Vol. 11424.
1386 Springer, 116–133. https://doi.org/10.1007/978-3-030-16722-6_7
- 1387 [14] Lars Fritzsche, Jens Kosiol, Andy Schürr, and Gabriele Taentzer.
1388 2020. Avoiding Unnecessary Information Loss: Correct and Efficient
1389 Model Synchronization Based on Triple Graph Grammars. *CoRR*
1390 [abs/2005.14510](https://arxiv.org/abs/abs/2005.14510) (2020). arXiv:2005.14510
- 1391 [15] Lars Fritzsche, Erhan Leblebici, Anthony Anjorin, and Andy Schürr.
1392 2017. A Look-Ahead Strategy for Rule-Based Model Transformations.
1393 In *Proceedings of MODELS 2017 Satellite Event: Workshops (ModComp,
1394 ME, EXE, COMMITMDE, MRT, MULTI, GEMOC, MoDeVVA, MDETools,
1395 FlexMDE, MDEbug), Posters, Doctoral Symposium, Educator Symposium,
1396 ACM Student Research Competition, and Tools and Demonstrations
1397 co-located with ACM/IEEE 20th International Conference on Model
1398 Driven Engineering Languages and Systems (MODELS 2017), Austin,
1399 TX, USA, September, 17, 2017 (CEUR Workshop Proceedings)*, Loli Bur-
1400 gueño, Jonathan Corley, Nelly Bencomo, Peter J. Clarke, Philippe Col-
1401 let, Michalis Famelis, Sudipto Ghosh, Martin Gogolla, Joel Greenyer,
1402 Esther Guerra, Sahar Kokaly, Alfonso Pierantonio, Julia Rubin, and
1403 Davide Di Ruscio (Eds.), Vol. 2019. CEUR-WS.org, 45–53.
- 1404 [16] Holger Giese and Robert Wagner. 2009. From model transformation to
1405 incremental bidirectional model synchronization. *Journal of Software
1406 and Systems Modeling* 8, 1 (2009), 21–43. <https://doi.org/10.1007/s10270-008-0089-9>
- 1406 [17] Susann Gottmann, Frank Hermann, Nico Nachtigall, Benjamin Braatz,
1407 Claudia Ermel, Hartmut Ehrig, and Thomas Engel. 2013. Correctness
1408 and Completeness of Generalised Concurrent Model Synchronisation
1409 Based on Triple Graph Grammars. In *Proceedings of the Second Work-
1410 shop on the Analysis of Model Transformations (AMT 2013), Miami, FL,
1411 USA, September 29, 2013 (CEUR Workshop Proceedings)*, Benoit Baudry,
1412 Jürgen Dingel, Levi Lucio, and Hans Vangheluwe (Eds.), Vol. 1077.
1413 CEUR-WS.org.
- 1414 [18] Reiko Heckel and Gabriele Taentzer. 2020. *Graph Transformation for
1415 Software Engineers - With Applications to Model-Based Development
1416 and Domain-Specific Language Engineering*. Springer. <https://doi.org/10.1007/978-3-030-43916-3>
- 1417 [19] Frank Hermann, Hartmut Ehrig, Claudia Ermel, and Fernando Orejas.
1418 2012. Concurrent Model Synchronization with Conflict Resolution
1419 Based on Triple Graph Grammars. In *Fundamental Approaches to Soft-
1420 ware Engineering - 15th International Conference, FASE 2012, Held as
1421 Part of the European Joint Conferences on Theory and Practice of Soft-
1422 ware, ETAPS 2012, Tallinn, Estonia, March 24 - April 1, 2012. Proceed-
1423 ings (Lecture Notes in Computer Science)*, Juan de Lara and Andrea Zisman
1424 (Eds.), Vol. 7212. Springer, 178–193. https://doi.org/10.1007/978-3-642-28872-2_13
- 1425 [20] Frank Hermann, Hartmut Ehrig, Ulrike Golas, and Fernando Orejas.
1426 2010. Efficient analysis and execution of correct and complete
1427 model transformations based on triple graph grammars. In *Proceed-
1428 ings of the First International Workshop on Model-Driven Interoper-
1429 ability, MDI@MoDELS 2010, Oslo, Norway, October 3-5, 2010*, Jean
1430 Bézivin, Richard Mark Soley, and Antonio Vallecillo (Eds.). ACM, 22–31.
1431 <https://doi.org/10.1145/1866272.1866277>
- 1432 [21] Frank Hermann, Hartmut Ehrig, Fernando Orejas, Krzysztof Czar-
1433 necki, Zinovy Diskin, Yingfei Xiong, Susann Gottmann, and Thomas
1434 Engel. 2015. Model synchronization based on triple graph grammars:
1435 correctness, completeness and invertibility. *Journal of Software and
1436 Systems Modeling* 14, 1 (2015), 241–269. <https://doi.org/10.1007/s10270-012-0309-1>

- 1431 [22] Thomas Hettel, Michael Lawley, and Kerry Raymond. 2008. Model
1432 Synchronisation: Definitions for Round-Trip Engineering. In *Theory
1433 and Practice of Model Transformations - 1st International Conference,
1434 ICMT@TOOLS 2008, Zurich, Switzerland, July 1-2, 2008, Proceedings
1435 (Lecture Notes in Computer Science)*, Antonio Vallecillo, Jeff Gray, and
1436 Alfonso Pierantonio (Eds.), Vol. 5063. Springer, 31–45. https://doi.org/10.1007/978-3-540-69927-9_3
- 1437 [23] Timo Kehrer. 2015. *Calculation and propagation of model changes based
1438 on user-level edit operations: a foundation for version and variant man-
1439 agement in model-driven engineering*. Ph.D. Dissertation. University of
Siegen.
- 1440 [24] Felix Klar, Marius Lauder, Alexander Königs, and Andy Schürr. 2010. Extended Triple Graph Grammars with Efficient and Compatible Graph Translators. In *Graph Transformations and Model-Driven Engineering - Essays Dedicated to Manfred Nagl on the Occasion of his 65th Birthday*, Gregor Engels, Claus Leverentz, Wilhelm Schäfer, Andy Schürr, and Bernhard Westfechtel (Eds.). Lecture Notes in Computer Science, Vol. 5765. Springer, 141–174. https://doi.org/10.1007/978-3-642-17322-6_8
- 1441 [25] Dimitrios S. Kolovos, Richard F. Paige, and Fiona Polack. 2008. Detecting and Repairing Inconsistencies across Heterogeneous Models. In *First International Conference on Software Testing, Verification, and Validation, ICST 2008, Lillehammer, Norway, April 9-11, 2008*. IEEE Computer Society, 356–364. <https://doi.org/10.1109/ICST.2008.23>
- 1442 [26] Jens Kosiol, Lars Fritzsche, Andy Schürr, and Gabriele Taentzer. 2020. Double-pushout-rewriting in S-Cartesian functor categories: Rewriting theory and application to partial triple graphs. *Journal of Logical and Algebraic Methods in Programming* 115 (2020), 100565.
- 1443 [27] Erhan Leblebici. 2018. *Inter-Model Consistency Checking and Restoration with Triple Graph Grammars*. Ph.D. Dissertation. Darmstadt University of Technology, Germany.
- 1444 [28] Nuno Macedo and Alcino Cunha. 2016. Least-change bidirectional model transformation with QVT-R and ATL. *Journal of Software and Systems Modeling* 15, 3 (2016), 783–810. <https://doi.org/10.1007/s10270-014-0437-x>
- 1445 [29] Fernando Orejas, Artur Boronat, Hartmut Ehrig, Frank Hermann, and Hanna Schölzel. 2013. On Propagation-Based Concurrent Model Synchronization. *Electronic Communication of the European Association of Software Science and Technology* 57 (2013). <https://doi.org/10.14279/tuj.eceast.57.871>
- 1446 [30] Fernando Orejas, Elvira Pino, and Marisa Navarro. 2020. *Incremental Concurrent Model Synchronization using Triple Graph Grammars*. Lecture Notes in Computer Science, Vol. 12076. Springer, 273–293. https://doi.org/10.1007/978-3-030-45234-6_14
- 1447 [31] Benjamin Pierce, A. Schmitt, and Michael Greenwald. 2003. Bringing harmony to optimism: a synchronization framework for heterogeneous tree-structured data. *Technical Report MS-CIS-03-42* (01 2003).
- 1448 [32] Andy Schürr. 1994. Specification of Graph Translators with Triple Graph Grammars. In *Graph-Theoretic Concepts in Computer Science, 20th International Workshop, WG '94, Herrsching, Germany, June 16-18, 1994, Proceedings (Lecture Notes in Computer Science)*, Ernst W. Mayr, Gunther Schmidt, and Gottfried Tinhofer (Eds.), Vol. 903. Springer, 151–163. https://doi.org/10.1007/3-540-59071-4_45
- 1449 [33] Gabriele Taentzer, Manuel Ohrndorf, Yngve Lamo, and Adrian Rutle. 2017. Change-Preserving Model Repair. In *Fundamental Approaches to Software Engineering - 20th International Conference, FASE 2017, Held as Part of the European Joint Conferences on Theory and Practice of Software, ETAPS 2017, Uppsala, Sweden, April 22-29, 2017, Proceedings (Lecture Notes in Computer Science)*, Marieke Huisman and Julia Rubin (Eds.), Vol. 10202. Springer, 283–299. https://doi.org/10.1007/978-3-662-54494-5_16
- 1450 [34] Laurence Tratt. 2008. A change propagating model transformation Language. *Journal of Object Technology* 7, 3 (2008), 107–124. <https://doi.org/10.5381/jot.2008.7.3.a3>
- 1451 [35] Frank Trollmann and Sahin Albayrak. 2017. Decision Points for Non-determinism in Concurrent Model Synchronization with Triple Graph Grammars. In *Theory and Practice of Model Transformation - 10th International Conference, ICMT@STAF 2017, Marburg, Germany, July 17-18, 2017, Proceedings (Lecture Notes in Computer Science)*, Esther Guerra and Mark van den Brand (Eds.), Vol. 10374. Springer, 35–50. https://doi.org/10.1007/978-3-319-61473-1_3
- 1452 [36] Nils Weidmann, Anthony Anjorin, Lars Fritzsche, Gergely Varró, Andy Schürr, and Erhan Leblebici. 2019. Incremental Bidirectional Model Transformation with eMoflon: iBeX. In *Proceedings of the 8th International Workshop on Bidirectional Transformations co-located with the Philadelphia Logic Week, Bx@PLW 2019, Philadelphia, PA, USA, June 4, 2019 (CEUR Workshop Proceedings)*, James Cheney and Hsiang-Shang Ko (Eds.), Vol. 2355. CEUR-WS.org, 45–55.
- 1453 [37] Yingfei Xiong, Zhenjiang Hu, Haiyan Zhao, Song Hui, Hong Mei, Yingfei Xiong, Haiyan Zhao, Zhenjiang Hu, Masato Takeichi, Song Hui, and Hong Mei. 2008. Beanbag: Operation-based Synchronization with IntraRelations. In *Grace Technical Reports, GRACE-TR-2008-04*. National Institute of Informatics.
- 1454 [38] Yingfei Xiong, Hui Song, Zhenjiang Hu, and Masato Takeichi. 2009. Supporting Parallel Updates with Bidirectional Model Transformations. In *Theory and Practice of Model Transformations - 2nd International Conference, ICMT@TOOLS 2009, Zurich, Switzerland, June 29-30, 2009. Proceedings (Lecture Notes in Computer Science)*, Richard F. Paige (Ed.), Vol. 5563. Springer, 213–228. https://doi.org/10.1007/978-3-642-02408-5_15
- 1455 [39] Yingfei Xiong, Hui Song, Zhenjiang Hu, and Masato Takeichi. 2013. Synchronizing concurrent model updates based on bidirectional transformation. *Journal of Software and Systems Modeling* 12, 1 (2013), 89–104. <https://doi.org/10.1007/s10270-010-0187-3>

A Overview over conflicts and actions for propagation

In Table 1 we give an overview of the different possible conflicts depending on the annotation of the DPC.

B Evaluation Measurements

Figure 19 and Fig. 20 depict the evaluation measurements of our second TGG, which is based on Modisco [5] and a custom documentation model. The description for these measurements can be found in Section 7.

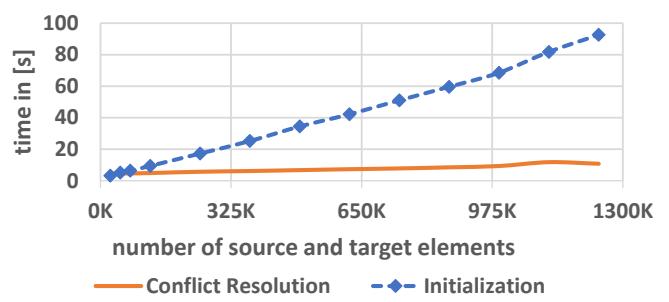


Figure 19. Increasing model size with constant number of 100 conflicts

1486
1487
1488
1489
1490
1491
1492
1493
1494
1495
1496
1497
1498
1499
1500
1501
1502
1503
1504
1505
1506
1507
1508
1509
1510
1511
1512
1513
1514
1515
1516
1517
1518
1519
1520
1521
1522
1523
1524
1525
1526
1527
1528
1529
1530
1531
1532
1533
1534
1535
1536
1537
1538
1539
1540

	target source	{}	-	/	#	n	
1541	{}		pdc	pdc, cpc	cpc		1596
1542	-		pdc	pdc, cpc	pdc	pdc, cpc	1597
1543	/	pdc, cpc	1598				
1544	#		pdc	pdc, cpc	acc	cpc	1599
1545	n	cpc	pdc, cpc	pdc, cpc	cpc	cpc	1600
1546							1601
1547							1602
1548							1603

Table 1. Overview of potential conflicts depending on the annotations in the DPG where *acc* stands for attribute change, *pdc* for preserve-deletion, and *cpc* for correspondence preservation conflict.

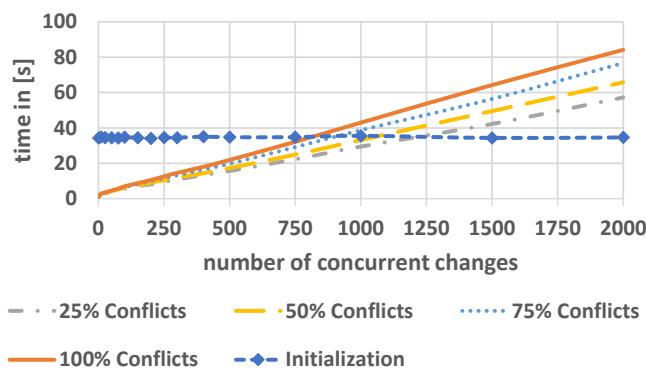


Figure 20. Increasing number of changes with constant model of 500K nodes

C Intermediate conflict resolution steps

Figure 21 depicts the model of our running example after *Local CC*, *Translate* and *Repair* have been applied, while Fig. 22 shows the corresponding delta precedence graph with conflicts \hat{C}_1 , \hat{C}_2 and \hat{C}_3 . Figure 23 and Fig. 24 depict the final model and delta precedence graph, respectively.

1596
1597
1598
1599
1600
1601
1602
1603
1604
1605
1606
1607
1608
1609
1610
1611
1612
1613
1614
1615
1616
1617
1618
1619
1620
1621
1622
1623
1624
1625
1626
1627
1628
1629
1630
1631
1632
1633
1634
1635
1636
1637
1638
1639
1640
1641
1642
1643
1644
1645
1646
1647
1648
1649
1650

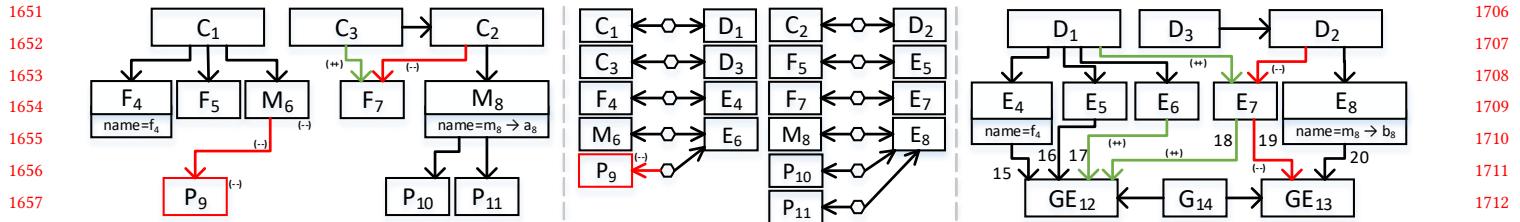


Figure 21. Running Example – Model and Delta after LocalCC, Repair and Translate

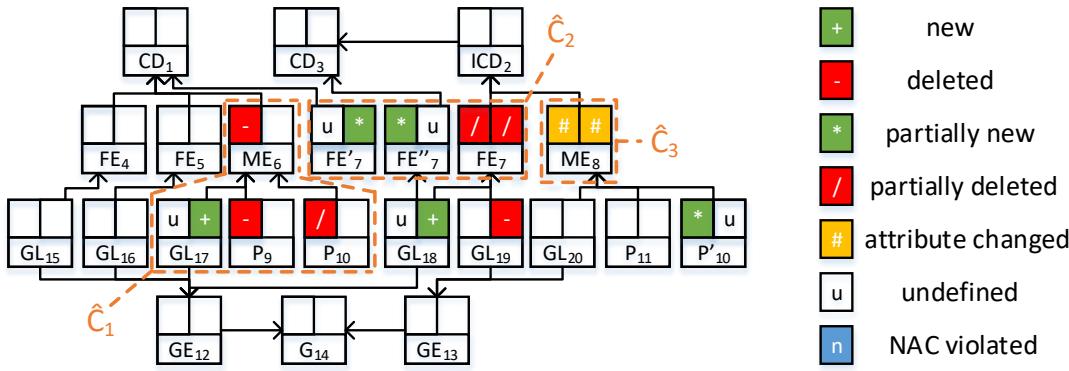


Figure 22. Running Example – Delta Precedence Graph after LocalCC, Repair and Translate

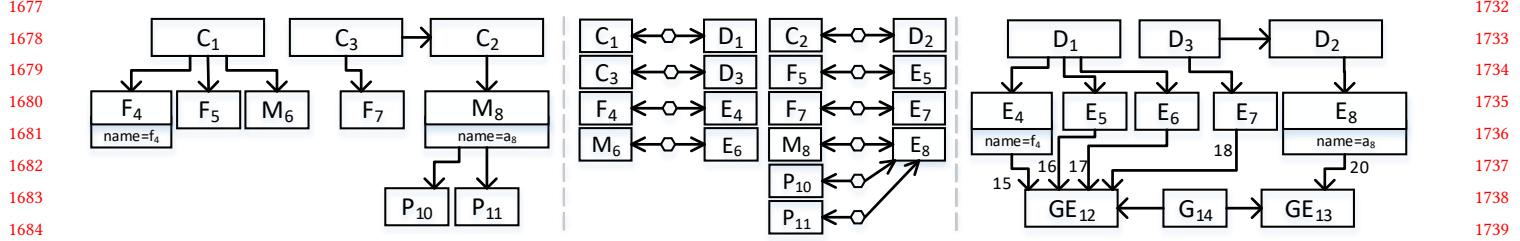


Figure 23. Running Example – Synchronized Model

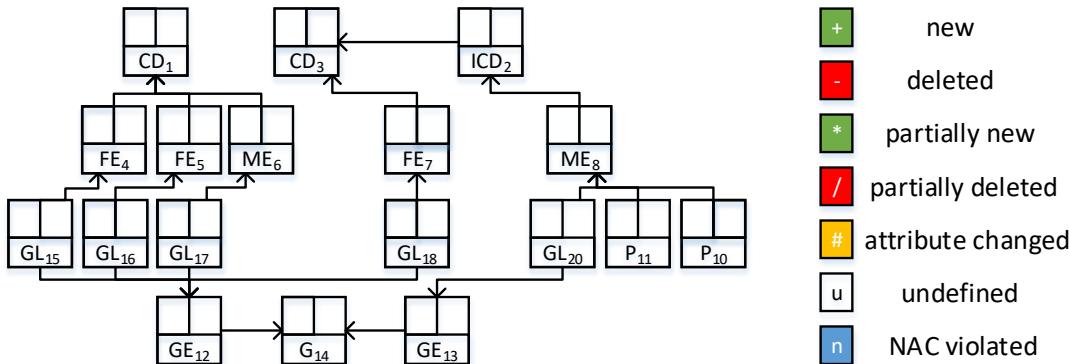


Figure 24. Running Example – Synchronized Precedence Graph

1761	1816
1762	1817
1763	1818
1764	1819
1765	1820
1766	1821
1767	1822
1768	1823
1769	1824
1770	1825
1771	1826
1772	1827
1773	1828
1774	1829
1775	1830
1776	1831
1777	1832
1778	1833
1779	1834
1780	1835
1781	1836
1782	1837
1783	1838
1784	1839
1785	1840
1786	1841
1787	1842
1788	1843
1789	1844
1790	1845
1791	1846
1792	1847
1793	1848
1794	1849
1795	1850
1796	1851
1797	1852
1798	1853
1799	1854
1800	1855
1801	1856
1802	1857
1803	1858
1804	1859
1805	1860
1806	1861
1807	1862
1808	1863
1809	1864
1810	1865
1811	1866
1812	1867
1813	1868
1814	1869
1815	1870